Open Retrieval Study

SYNTHETIC EDUCATIONAL FIXTURE - NOT A PUBLISHED RESEARCH PAPER

Abstract: An example document for testing page-grounded retrieval.

Research question: How can a lightweight assistant retrieve auditable evidence?

Method: Character n-grams are mapped to sparse hashing vectors.

Cosine similarity ranks the chunks without a GPU or embedding API.

Contributions here are illustrative; this document claims no scientific novelty.

Experiments and Reproducibility

Dataset: 12 hand-written educational sentences.

The example optimizer is Adam with learning rate 0.001 and batch size 16.

This is a retrieval test fact, not an actual model training result.

Environment: Python 3.11 or later; CPU execution is sufficient.

Limitations: Lexical retrieval has no cross-language semantic guarantee.

No accuracy score, timing benchmark or unperformed experiment is claimed.